

# Bandwidth-Sensitive Oblivious Routing

by

Tina Wen

Submitted to the Department of Electrical Engineering and Computer  
Science

in partial fulfillment of the requirements for the degree of  
Master of Engineering in Computer Science and Engineering

at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

May 2009

© Massachusetts Institute of Technology 2009. All rights reserved.

Author .....  
Department of Electrical Engineering and Computer Science  
May 12, 2009

Certified by .....  
Srinivas Devadas  
Associate Department Head, Professor  
Thesis Supervisor

Accepted by .....  
Arthur C. Smith  
Professor of Electrical Engineering  
Chairman, Department Committee on Graduate Students



# Bandwidth-Sensitive Oblivious Routing

by

Tina Wen

Submitted to the Department of Electrical Engineering and Computer Science  
on May 12, 2009, in partial fulfillment of the  
requirements for the degree of  
Master of Engineering in Computer Science and Engineering

## Abstract

Traditional oblivious routing algorithms either do not take into account the bandwidth demand, or assume that each flow has its own private channel to guarantee deadlock freedom. Though adaptive routings can adapt to the network traffic, they require complicated router designs. In this thesis, we present a polynomial time heuristic routing algorithm that takes bandwidth requirements of each flow into account to minimize maximum channel load. On top of the heuristic algorithm, we have two variations. The first one produces a deadlock free route. The second one produces a minimal route, and is deadlock free with more than 2 virtual channels with proper VC allocation. Both routing algorithms are oblivious, and need only simple router designs. The performance of each bandwidth-sensitive routing algorithm is evaluated against dimension-order routing and against each other on a number of benchmarks.

Thesis Supervisor: Srinivas Devadas

Title: Associate Department Head, Professor



# Acknowledgments

I would like to dedicate this work to my advisor, Professor Srinivas Devadas. Srini is truly the best thesis advisor one can ever have. It has been an enjoyable learning journey for me throughout the two terms spent working on my thesis, and it would not be possible without the guidance and support Professor Devadas provided me.

I also would like to express my sincere gratitude to the four great people who I have been working with for the past year, Myong Hyon (Brandon) Cho, Michel Kinsy, Keun Sup Shim, and Mieszko Lis. I started working on my thesis in the middle of the adaptive routing project. Brandon and Michel gave me tremendous help to get me up to speed with the project. It would have taken me a lot more time and effort if they did not help me as much as they did. I greatly appreciate these two more experienced people in the group. Group discussion is the place where I learn the most. I enjoy every group meeting when we can exchange ideas, learn from each other, and explore new areas. Having Srini and these creative people on board, I learnt so many things that I could never have learnt from reading books or taking classes.

Brandon Cho, Keun Sup Shim, Michel Kinsy, and I also sat in the same office. We are not only work buddies, but also very good friends. I appreciate all the help and support they gave me in the past year. It was the sweetest thing in the world that Brandon and Keun Sup got me a birthday cake on my birthday and sang me happy birthday for me. With work buddies like this, how can I not enjoy my work?

I would also like to thank my boyfriend, Ramsey Khalaf, who I discussed some of my thesis work with from time to time. He helped me brain-storm and came up with many great ideas for my thesis project. He also helped me revising my thesis, which has been tremendous and valuable help for me from a native English speaker. I also want to thank him for all the support he has given me. He is always there when I am stressed or confused.

I also greatly appreciate some of my very good friends who helped my thesis in some way or another. Victor Costan, a friend who I knew from way back when he was my 6.006 TA, provided a lot of academic and mental support for my thesis. Danh

Vo, the smartest hardware student on earth I know, gave me a lot of support and hardware advice.

I would like to thank my family and friends who contributed in making me who I am today. I am truly glad that I grew up in Beijing, China. The eighteen years experienced living in China made me a strong, dedicated, and persistent girl. I would not be able to finish this thesis project if I was not so dedicated.

Finally I'd like to thank MIT for accepting me as a student. These past five years (four-year undergraduate and one-year MEng) have been a life-changing experience for me. I not only grew to be a good programmer and electrical engineer, but also became an interesting and pleasant person to be around with. MIT contributed significantly in making who I am today, and I enjoyed every bit of it.

# Contents

<b>1</b>	<b>Introduction</b>	<b>15</b>
1.1	Thesis Outline . . . . .	17
<b>2</b>	<b>Background and Related Work</b>	<b>19</b>
2.1	Diastolic Arrays . . . . .	20
2.2	Oblivious and Adaptive Routing . . . . .	21
2.2.1	XY/YX Routing . . . . .	21
2.2.2	ROMM and Valiant Routing . . . . .	22
2.3	Router Designs . . . . .	23
2.3.1	Typical Virtual Channel Router . . . . .	23
2.3.2	Source Routing and Node-Table Routing . . . . .	23
2.3.3	Virtual Channels . . . . .	24
2.4	Linear Programming . . . . .	25
<b>3</b>	<b>Routing using Dijkstra's algorithm</b>	<b>27</b>
3.1	Dijkstra's Routing Algorithm . . . . .	28
3.2	Setting the Weights . . . . .	29
3.3	Breaking Cycles using the Turn Model . . . . .	32
3.4	Nodes With 4 States . . . . .	34
3.5	Modified Dijkstra: 1 Iteration vs. 100 Iterations . . . . .	36
3.6	Output Selection Process . . . . .	39
<b>4</b>	<b>Experiment and Performance: Dijkstra's Routing Algorithm</b>	<b>41</b>

4.1	Simulation Process . . . . .	42
4.1.1	The Simulator . . . . .	42
4.1.2	Benchmarks . . . . .	43
4.2	Performance Evaluation of Dijkstra's Algorithm . . . . .	44
4.2.1	Channel Load Comparison . . . . .	44
4.2.2	Simulation Results of 4 Synthetic Benchmarks . . . . .	45
4.2.3	Simulation Result with Bandwidth Variations . . . . .	47
<b>5</b>	<b>Dijkstra's Routing Algorithm in Minimal Routing</b>	<b>51</b>
5.1	Motivations for exploring Dijkstra's Algorithm in Minimal Routing .	52
5.2	Virtual Channel Allocation for Dijkstra's Algorithm in Minimal Routing	53
5.3	Implementation of Bandwidth-Sensitive Minimal Routing using Dijkstra's Algorithm . . . . .	55
5.3.1	Elimination of Turn Model Constraint . . . . .	55
5.3.2	Elimination of Four States Per Node . . . . .	55
5.3.3	The Addition of Making Some Progress at Every Step . . . . .	56
5.3.4	The Need of 100 Iterations . . . . .	56
5.4	Dijkstra's Routing Algorithm in Minimal Routing with Favored Turns	58
5.4.1	Implementation For Favoring Four Turns . . . . .	59
5.4.2	Advantage of this Implementation . . . . .	60
5.5	Four Turn Model Pairs for Minimal Routing . . . . .	62
5.6	Routes Selection for Dijkstra's Algorithm in Minimal Routing . . . .	63
<b>6</b>	<b>Performance Comparison Between Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing</b>	<b>65</b>
6.1	Channel Load Comparison . . . . .	66
6.2	Performance Comparison of 2 VCs on Four Benchmarks . . . . .	66
6.3	Performance Comparison of multiple VCs . . . . .	70
6.3.1	Shuffle with 2 and 4 VCs . . . . .	70
6.3.2	H.264 with private channels . . . . .	71



<b>7 Conclusion and Future Work</b>	<b>73</b>
<b>A Acronyms</b>	<b>75</b>



# List of Figures

3-1	Two weight functions over residual capacity . . . . .	30
3-2	A situation where the current Dijkstra's Routing Algorithm would not produce the correct answer. . . . .	34
3-3	Routing result for routing S to D1 first and S to D2 second . . . . .	36
3-4	Routing result for routing S to D2 first . . . . .	37
3-5	Head-of-line blocking illustration. There is heavy congestion ahead of flow A, and no congestion ahead of flow B after the split. . . . .	40
4-1	Routes generated using Dijkstra's Routing Algorithm for Bit-complement.	45
4-2	Routes generated using Dijkstra's Routing Algorithm for Transpose. .	46
4-3	Load-throughput graphs for bit-complement on a router with 1 or 2 VCs. . . . .	48
4-4	Load-throughput graphs for transpose on a router with 1 or 2 VCs. .	48
4-5	Load-throughput graphs for shuffle on a router with 1 or 2 VCs . . . .	49
4-6	Load-throughput graphs for H.264 on a router with 1 or 2 VCs. . . .	49
4-7	Load-throughput graphs for transpose (1 virtual channel) when band- widths change by $\pm 10\%$ and $\pm 50\%$ after route computation. . . . .	50
5-1	The West-First Turn Model. . . . .	55
5-2	The North-Last Turn Model. . . . .	56
5-3	4 possible minimal routes from S to D where S and D are 2 units in length and 2 units in width apart from each other. . . . .	57
5-4	The four (out of possible eight) different one-turn routes on a 2-dimensional mesh that conform to both the West-First and North-Last Turn Model.	58

5-5	Two routes to illustrate favored routes. The solid route is a favored route while the dotted route is not. . . . .	59
5-6	West-First and North-Last Turn Models rotated four times to form 4 Turn Model pairs. The column WF shows the West-First Turn Model rotated 4 times. The column NL shows the North-Last Turn Model rotated 4 times. The last column shows the favored minimal routes that can be assigned to either virtual channel. . . . .	62
6-1	Performance comparison for bit-complement on a router with 2 VCs.	67
6-2	Performance comparison for transpose on a router with 2 VCs. . . . .	67
6-3	Performance comparison for shuffle on a router with 2 VCs. . . . .	68
6-4	Performance comparison for H.264 on a router with 2 VCs. . . . .	68
6-5	Performance comparison for Shuffle on a router with 2 and 4 VCs. . .	70
6-6	Performance comparison for H.264 on a router with 10 VCs. . . . .	72

# List of Tables

3.1	Turns prohibited for 20 cases in the Turn Model . . . . .	33
4.1	Comparison of Maximum Channel Load (MCL) in MB/second. . . .	44
6.1	Comparison of MCL in MB/second for two variations of routing algorithms based on Dijkstra's Algorithm. . . . .	66



# Chapter 1

## Introduction

Routing methods can be divided into two main types, oblivious and adaptive [19]. In an oblivious routing algorithm a route is completely determined given a set of source and destination positions. Oblivious routing methods need very simple router designs. A lookup table is enough. No adapt capability in the routers is needed. However, most existing oblivious routing algorithms do not perform well, especially under conditions where an application has different bandwidth requirements for each flow. Adaptive routing, on the other hand, can adapt based on network traffic information, such as congestion. The network forms a feedback loop so it can adjust itself. However, adaptive routing requires complex router designs to make intelligent decisions and so adds extra hardware and delay.

Many techniques have been explored in designing Network-on-Chip (NoC) interconnect. [1] presents a survey. A large body of work considers routing in the NoC design phase when mapping applications onto NoC architectures [13], [17], [11]. This thesis adopts the idea of routing in the NoC design phase but it is unique in the sense that it iteratively uses a heuristic function to find the globally optimal routes.

A heuristic algorithm is introduced in [27] to improve the initial set of routes given by dimension-order routing and guarantees a deadlock free route. Palesi [20] developed a deadlock free, bandwidth-aware, adaptive routing algorithm for a specific application. Load is uniformly balanced across the network, so the routes globally perform well.

In this thesis, we propose and evaluate two bandwidth-sensitive oblivious routing algorithms that run in polynomial time and statically produce a set of deadlock free routes. Assuming knowledge of rough estimates of bandwidth demand for each flow, we heuristically minimize congestion at each link while satisfying each flow's demand and minimize the maximum channel load of the network. Bandwidth requirements are acquired by application analysis or profiling. After performing off-line routing and virtual channel allocation, next hop and virtual channel numbers corresponding to the flows are statically loaded onto the network while the processing elements are configured for computation. These two bandwidth-sensitive oblivious routing algorithms give higher throughput than traditional oblivious routing algorithms because the routes are optimized based on the global knowledge of bandwidth demand. Furthermore, the routers needed by these oblivious routing algorithms are simple and require hardly any hardware modification.

Both of the routing algorithms are based on Dijkstra's algorithm to find shortest paths. They both search iteratively through flows to balance bandwidth loads and to reduce maximum channel load. The first routing algorithm, Dijkstra's Routing Algorithm, produces a set of deadlock free routes that fit into a particular Turn Model. The second routing algorithm, Dijkstra's Algorithm in Minimal Routing, produces a set of minimal routes that are not deadlock free. However, they can be made deadlock free after virtual channel allocation with at least two virtual channels. Virtual channel allocation is done to firstly make the routes deadlock free (for the case of Dijkstra's Algorithm in Minimal Routing), secondly to reduce head-of-line blocking, and thirdly to minimize the number of flows per virtual channel.

The performance of routes produced by the two bandwidth-sensitive oblivious routing algorithms and routes produced by dimension-order routing, XY/YX, are compared on four benchmarks.



## 1.1 Thesis Outline

This thesis presents two methods to generate deadlock free bandwidth sensitive oblivious routes. It details the implementation of each and compares their performances against XY/YX routings as well as each other.

The rest of this thesis is structured as follows. Chapter 2 summarizes background and related work in applications, routing, and router design. Chapter 3 details every design issue Dijkstra's Routing Algorithm faces building on Dijkstra's Algorithm to find shortest paths. Chapter 4 describes the simulation process and compares maximum channel load and the performance results of Dijkstra's Routing Algorithm with Dimension-Order Routing, XY and YX, on four benchmarks. The effect of bandwidth variation is also simulated and analyzed. Chapter 5 details the implementation of Dijkstra's Algorithm in Minimal Routing. It is based on Dijkstra's Routing Algorithm and so the differences are compared. Chapter 6 compares the performance of Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing using 2, 4, and 8 virtual channels. Chapter 7 concludes this thesis and discusses future work.



# Chapter 2

## Background and Related Work

Before diving into the detailed implementation and performance analysis of our routing algorithms, this chapter talks about the basic structure of diastolic arrays in Section 2.1. Dijkstra’s Routing Algorithm, introduced in Chapter 3, is suitable for routing on these arrays. This chapter also discusses the differences and trade-offs between oblivious and adaptive routing methods in Section 2.2. Section 2.2.1 gives one example of oblivious routes, produced by the XY/YX routing algorithm. Section 2.2.2 discusses two other oblivious routing algorithms, ROMM and Valiant. Section 2.3 talks about some well-known simple router designs that are suitable for oblivious routing algorithms.

## 2.1 Diastolic Arrays

In [2] Cho proposed a diastolic array structure in which the processing elements communicate exclusively through First-In First-Out queues. Diastolic arrays are targeted at throughput sensitive, latency tolerant applications. Usually applications running on diastolic arrays are spatially partitioned into processing elements (PEs) and the network traffic pattern between these PEs remains relatively constant. This is because each PE performs a specific fixed task. The bandwidth demand of each flow stays roughly the same throughout. This thesis explores routing for a network whose traffic pattern is relatively static, rather than one with dramatic fluctuations in bandwidth demand. This is the case for most applications in reconfigurable computing. These applications are suitable for running on diastolic arrays.

## 2.2 Oblivious and Adaptive Routing

There are two main types of routing strategies, adaptive and oblivious, which are introduced in [19]. In adaptive routing, the path for each packet in a flow is dynamically determined by some global or local metric, such as the network traffic load. In oblivious routing, a packet's path is routed statically beforehand. Each flow defined by a source-destination pair always takes a predetermined path. In other words, the path in oblivious routing is completely determined by the source and destination positions.

Deterministic routing is a subset of oblivious routing. Given the source and destination positions, there is only one possible path and it is always chosen.

On the other hand, in adaptive routing routes may be changed during execution based on some global or local network information. This may imply better performance because of its ability to adapt to the network. However, router complexity increases to deal with the workload for collecting and processing network data.

### 2.2.1 XY/YX Routing

DOR (dimension order routing)[4] becomes XY and YX on a 2-D mesh. DOR is deterministic. XY routing chooses routes with the shortest possible distance from source to destination, travelling along the x direction first, then along the y direction. YX routing routes along the y direction first, then along the x direction. Given a source-destination pair, deterministic routing always chooses the same path [3]. Since there is no one single method that always performs best for all patterns of network traffic, XY/YX may not be the best routing strategy due to its determinism.

Although deterministic, XY/YX routing guarantees that the routes it produces are minimal; they have the fewest number of hops possible. This may not always be beneficial as it doesn't allow for the possibility of routing around a heavily congested link while its surrounding links are unused. In this situation a non-minimal route is likely to give better performance.

XY/YX is deterministic and minimal, which makes it impossible to take network

traffic information into account. It may perform very poorly under certain traffic patterns. However, XY/YX routing is deadlock free, so it provides us an easy solution that just works. Assuming no false resource dependencies, necessary and sufficient conditions for deadlock-free deterministic routing are given in [4]. We use this condition to determine if a set of routes is deadlock-free in our oblivious routing scheme introduced in Chapter 3.

### **2.2.2 ROMM and Valiant Routing**

ROMM [18] and Valiant [24] are two examples of oblivious routing methods. Dimension order routings are used by ROMM and Valiant. A middle node is randomly chosen to randomize the traffic pattern in order to distribute load better. In return [22], a worst-case throughput is guaranteed to be found by balancing traffic patterns in XY and YX. Weighted ordered toggle (WOT) [9] uses more than 2 virtual channels to reduce maximum network load for a given source to destination pair. These routes either do not take into account bandwidth demands, or are limited to simple minimal paths (WOT). Our goal is to develop a routing algorithm that is bandwidth sensitive, but also utilizes non-minimal routes.

## 2.3 Router Designs

### 2.3.1 Typical Virtual Channel Router

The architecture and operation of a typical virtual channel router can be found in [5], [21], [16]. The data path of a router consists of buffers and a switch with usually multiple buffers for each physical link. This way it appears to have multiple virtual channels available for transmission. When a packet is ready to be sent, the switch connects an input buffer to a corresponding output channel. There are three major control modules for managing the data path: a router, a virtual-channel (VC) allocator, and a switch allocator. These control modules determine the next hop, the next virtual channel and when a switch is available for the next flit, respectively.

There are four stages for a routing operation: routing (RC), virtual-channel allocation (VA), switch allocation (SA), and switch traversal (ST). These four stages are often represented as four pipeline stages in modern virtual channel routers. When a head flit (the beginning of a packet) arrives at an input channel, the next hop is determined by the routing (RC) stage. Based on the next hop, a virtual channel is allocated by the VA stage. Then the head flit needs to compete for a switch (SA). If the next hop is able to accept a flit, the flit moves on to the output port (ST).

Traditional oblivious routing algorithms such as Dimension Ordered Routing (DOR) [4], ROMM [18], Valiant [24], and o1turn [22], have next hops completely determined given a source destination pair. This allows the next hop to be computed quite easily at each router node using the packet's destination.

Oblivious routing algorithms need only simple router designs. Simple logic is needed for the routing (RC) stage. The clock frequency that the router runs on is usually determined by the switch allocation (SA) stage [21].

### 2.3.2 Source Routing and Node-Table Routing

Source routing and node-table routing are suitable for any static routes, regardless of whether the routes are minimal or not. Routers need to be programmable to

remember the information of the next hop, and flexible enough to support any fixed route from any source to any destination. Source routing and node-table routing cannot support adaptive routing algorithms.

Source routing and node-table routing are described in detail in [14]. Source routing creates longer packets as the next hop directions are carried in each packet. Node-table routing requires a lookup table in each node to store which direction a packet should be sent as well slightly larger packets to store the table index. In either case, there is no additional computation logic in comparison with the router described in 2.3.1. In fact, source routing and node-table routing are suitable for any oblivious routing algorithm.

### **2.3.3 Virtual Channels**

There are usually many virtual channels in modern routers, which may support dynamic virtual channel allocation. In the single virtual channel case, cycles in cycle dependency graph imply the network may deadlock. It is possible to have a deadlock free path with a cycle in cycle dependency graph if an escape path is provided [7], [8]. However, this involves adaptive routing, which we do not consider in this thesis.



## 2.4 Linear Programming

A Linear Programming formulation can be used to determine a lower bound on maximum channel load given a set of flows. However, the routes given by linear programming may not be realizable by modern routers, because a packet flow may be split into different paths (disjoint flow) to achieve maximum throughput. This may result in packets received out of order. Besides, the routes produced may not be deadlock free.

A route in which every packet flows along a single path is called a **joint flow**. However, the joint flow problem is NP-hard even for single sources [15], requiring the use of estimation algorithms or heuristics for large problem sets. Mixed integer-linear programming (MILP) can produce an optimal result by minimizing maximum channel load (MCL), for small problems. MILP is NP-hard but produces routes with joint flows. However, a route with the lowest possible MCL may not be the optimal route for performance. MILP motivates us to find a polynomial time algorithm that can produce an optimal or close to optimal set of routes with joint flows.



## Chapter 3

# Routing using Dijkstra's algorithm

XY and YX routings are minimal and deterministic, but are purely based on the source and destination positions and do not take into account network traffic information. If we could profile the applications before running and record bandwidth requirement for each source-destination pair, we could utilize this information in our routing algorithm, and thus generate better routes. This chapter proposes an oblivious routing algorithm based on Dijkstra's Algorithm for finding shortest paths. It provides a polynomial time solution to the routing problem given knowledge of each flows bandwidth requirement. Since the routing algorithm is oblivious, very little hardware support is needed from the routers. Simple router designs that support source routing or node-table routing will do.

Traditionally Dijkstra's Algorithm is used to find the shortest paths from one source to all destinations on a directed graph. We can model the processing elements in diastolic arrays in [2] as nodes and the connections between processing elements as edges on our directed graph. If we set the weights on the edges to reflect the residual capacity of each link, we can then run Dijkstra's Algorithm to find the paths with the lowest weight. This way, congested links are avoided and the resulting graph would have many links with large residual capacities, and therefore is likely to perform well under high traffic loads. This chapter describes the implementation of the oblivious routing algorithm.

### 3.1 Dijkstra's Routing Algorithm

Dijkstra's algorithm for finding shortest paths is described in [6]. It solves the single-source shortest path problem on a weighted, directed graph, where all edge weights are non-negative.

Ideally we would like to have routes that avoid using links with little residual capacity, so the network does not get congested. We can incorporate this idea by using residual capacity information to modify the weights on the graph's edges. When the residual capacity is high, the weight is small and when a link is heavily used, its weight is increased. This modified Dijkstra's algorithm to find shortest paths for routing is named Dijkstra's Routing Algorithm. Dijkstra's Routing Algorithm is less liable to choose paths utilizing congested links as these edges have higher weights. The idea is that Dijkstra's Routing Algorithm can give us a route with low congestion by taking into account of bandwidth demands of each source-destination pair, and will therefore perform better than the deterministic routing algorithms, XY and YX.

The complexity of the algorithm depends upon the implementation of the heap. Here we chose a Fibonacci heap over an array or a binary min-heap for performance reasons. If  $V$  is the number of vertices and  $E$  is the number of edges in the graph, the complexity of `extract_min` for a Fibonacci heap is  $\Theta(\log V)$ , and the complexity of `decrease_key` is  $\Theta(1)$ . Therefore, the complexity of Dijkstra's algorithm before any modifications is  $\Theta(V \times \text{extract\_min} + E \times \text{decrease\_key}) = \Theta(V \log V + E)$ .

## 3.2 Setting the Weights

Dijkstra's algorithm is able to find paths with minimum weights from a source to all destinations on the graph, given that all weights are non-negative. In order to incorporate Dijkstra's algorithm into our bandwidth sensitive routing, we have to modify the weights to not only take into account of the length of the path, but also its bandwidth demand.

We run Dijkstra's algorithm on a weighted, directed graph, deriving the weights from the residual capacities of each link. Consider a link  $e$  that has a capacity  $c(e)$ . We create a variable for the link  $\tilde{c}(e)$  which is the current residual capacity of link  $e$ . Initially, it is equal to the capacity  $c(e)$  set to be a constant  $C$ . If a flow  $i$  is routed through this link  $e$ , we will subtract the demand  $d_i$  from the residual capacity. The residual capacity is always checked to see whether it is enough to supply the demand for the flow during routing. If there is not enough capacity, then Dijkstra's Routing Algorithm sets the weight of the link to be extremely high, so it never chooses the link. Therefore, the residual capacity  $\tilde{c}(e)$  will never be negative.

We have experimented with multiple ways of setting the weights. To take bandwidth into account, we decided to set the weight to be the reciprocal of link residual capacity, similar to the CSPF metric described in [25]. Our weight function is  $w(e) = \frac{K}{\tilde{c}(e)-d_i}$  except if  $\tilde{c}(e) \leq d_i$ , then  $w(e) = \infty$ , and the algorithm never chooses the link. The constant  $C$  is set to be the smallest number that can provide us with routes for all flows without using  $\infty$ -weight links. The maximum channel load (MCL)<sup>1</sup> from XY or YX routing gives us an upper bound for  $C$ , but in most cases, we can set  $C$  lower and still find a solution. We want to set  $C$  lower because a lower  $C$  makes the algorithm more aggressively avoid congested links due to their higher weight. As it is shown in Figure 3-1. The weight function indicated by solid line with circles has a lower  $C$  than the weight function indicated by solid line. We can see that links with lower residual capacity are punished even more on the solid line with circles as more effort is put into avoiding congested links. The Mixed Integer-Linear Programming

---

<sup>1</sup>MCL stands for Maximum Channel Load. This number is the maximum of link load for all links.

(MILP) approach for routing introduced in [26] generates routes by setting the joint flow constraint and minimizing MCL. The result given by MILP sets a lower bound for  $C$ , but the routes do not necessarily perform the best.  $C$  needs to be tuned along with  $K$  to get good results.  $K$  is a constant in the weight function. The higher we set  $K$ , the more effort put into minimizing the hop count. As shown in Figure 3-1, the dotted line shows a  $K$  value that is 3 times the original  $K$  shown by the solid line. The weight increases for every case of residual capacity  $\tilde{c}(e)$ . By setting higher weights corresponding to each  $\tilde{c}(e)$ , we effectively reduce hop counts and get shorter paths.

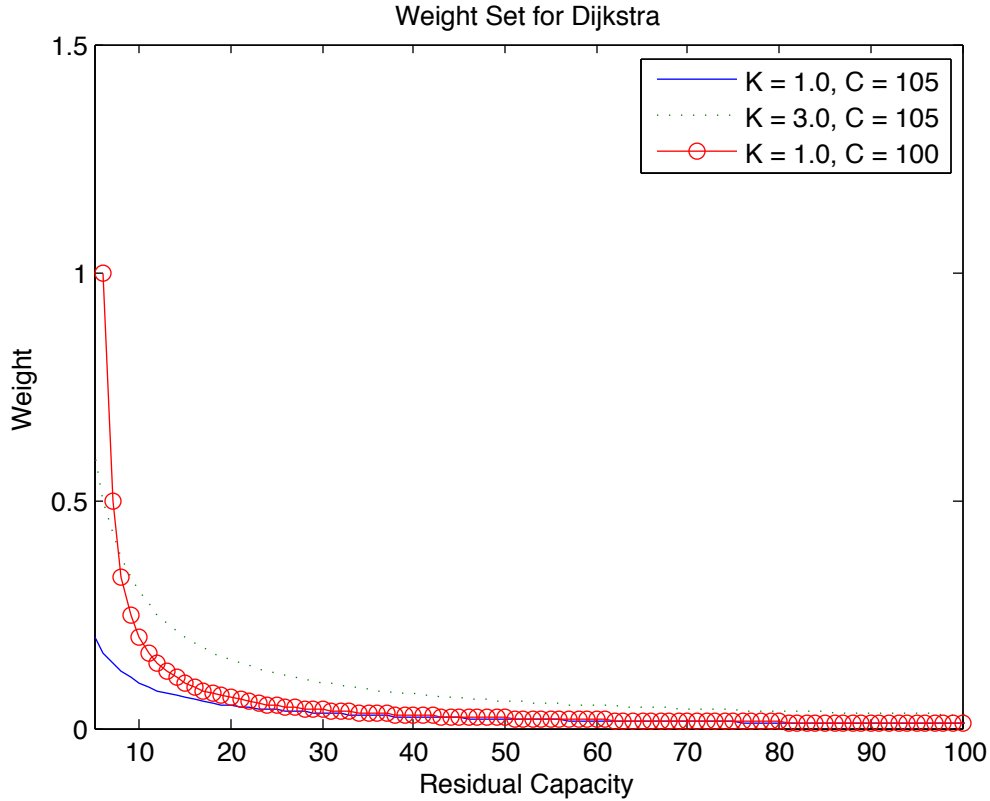


Figure 3-1: Two weight functions over residual capacity

From the above paragraph, we conclude that by tweaking  $C$  and  $K$  in the weight function, we set a trade-off between path length and link load. However, it is not the case that the more weight we give to the more congested links, the better. If we punish the congested links too much by setting  $C$  lower, the weights get extremely high for links with low residual capacity. Figure 3-1 shows the solid line with circles

having much higher weights than the solid line for low residual capacities. If this solid line with circles gets too high for low  $\tilde{c}(e)$ , the weights are too high for congested links then they matter so much that a slightly less congested links do not matter any more. In this situation, we may still end up getting many heavily congested links, however, these links are just a little less congested than the most congested ones, but they are still heavily congested. Caution needs to be taken when setting  $C$  and  $K$  in order to get an optimal solution.

### 3.3 Breaking Cycles using the Turn Model

Dijkstra’s Routing Algorithm guarantees us the shortest weight path for each source to destination pair. However, it doesn’t guarantee us a cycle free route. Cyclic routes may result in deadlock assuming limited buffer space.

In a 2-D mesh, there are 8 possible 90-degree turns. XY and YX routes do not have deadlock problems because each of them only uses 4 of the 8 90-degree turns. Furthermore, the 4 turns used are such that there is no way a cycle can be formed using them. However, XY and YX may be too restrictive for cycle breaking purposes.

The Turn Model introduced in [10] states that need only eliminate 2 of the 8 90-degree turns to eliminate cycles. If we assume that the mesh does not have wraparound channels and 0-degree turns, then we can also allow 2 out of the 4 180-degree turns. These turns to be eliminated are not random. As long as we follow the rules described in the Turn Model, we are guaranteed to get an acyclic set of routes which are thus deadlock free. Enumeration of the Turn Model is listed in Table 3.1.

Allowing 2 out of 4 180-degree turns is merely for completeness of the Turn Model. Incorporating 180-degree turns does not produce better routes. If a 180-degree turn is chosen, it only adds extra path length. Since Dijkstra’s Routing Algorithm tries to reduce each link’s weight, which takes into account of hop count, 180-degree turns never appear in the routes produced.

There are two ways to generate acyclic routes. The first is to generate an acyclic Cycle Dependency Graph (CDG) using the Turn Model, and then run the routing algorithm on it. Generating an acyclic CDG is described in [14]. The second way is to run Dijkstra’s Routing Algorithm straight on the mesh. However, during execution we keep track of the direction of the incoming edge, aka, the parent of the node, and so we know which turn is being made for each outgoing edge. If that turn is not allowed in a specific Turn Model, then we set the weight on the illegal output edge to be very high; higher than any valid route. In doing this we make sure Dijkstra’s Routing Algorithm will never choose to relax this edge and never choose this turn.

The complexity of the algorithm should not change after adding the Turn Model



case	2 prohibited 90-degree turns		2 prohibited 180-degree turns	
1	east-north	west-north	south-north	west-east
2	east-north	west-north	south-north	east-west
3	north-west	north-east	north-south	west-east
4	north-west	north-east	north-south	east-west
5	east-south	west-south	north-south	west-east
6	east-south	west-south	north-south	east-west
7	south-west	south-east	south-north	east-west
8	south-west	south-east	south-north	west-east
9	north-west	south-west	east-west	north-south
10	north-west	south-west	east-west	south-north
11	west-south	west-north	west-east	south-north
12	west-south	west-north	west-east	north-south
13	north-east	south-east	west-east	south-north
14	north-east	south-east	west-east	north-south
15	east-north	east-south	east-west	north-south
16	east-north	east-south	east-west	south-north
17	west-north	south-east	south-north	west-east
18	west-south	north-east	north-south	west-east
19	north-west	east-south	east-west	north-south
20	south-west	east-north	south-north	east-west

Table 3.1: Turns prohibited for 20 cases in the Turn Model

check. Parents of nodes are already tracked in the original Dijkstra’s algorithm for finding shortest paths. Comparing the relative positions of nodes to determine whether they satisfy a particular Turn Model is an  $O(1)$  computation. Therefore, the order of growth of our modified algorithm remains the same:  $\Theta(V \log V + E)$ , still polynomial.

This way, we generate a route that is deadlock free.

### 3.4 Nodes With 4 States

Dijkstra's Algorithm guarantees finding the shortest path from source to a particular node after it extracts the node from the heap and finishes the **Relax** step. After it is done with one node, it never goes back to relax the same node again. However, after incorporating the Turn Model constraint, the current Dijkstra's Routing Algorithm may not produce the correct shortest path, or a path at all when it should.




Figure 3-2: A situation where the current Dijkstra's Routing Algorithm would not produce the correct answer.

Consider the situation shown in Figure 3-2. We are trying to route from source **S** to destination **D**. At node **A**, our current algorithm would pick the solid path versus the dotted path, as it has the shortest distance of 3, versus the dotted path's distance of 50. However, if we are using Turn Model case 15 or 16 shown in Table 3.1, east-north, east-south, and east-west turns are not allowed. Therefore, the solid path cannot go south at node **A** because it cannot take an east-south turn. It can only go east. After choosing the link with load 1 after node **A**, the solid path still cannot go south because the east-south turn is prohibited. It can't go west either, because the east-west turn is prohibited. Therefore, the solid path is a dead-end and can never reach node **D**.

Our current Dijkstra's Routing Algorithm would not produce a valid route. However, in this case, the dotted path from **S** to **A** to **D** should have been chosen. It is a valid route with weight  $50 + 1 = 51$ .

One way to solve this problem is to create four nodes for each original node in our Dijkstra's Routing Algorithm. These new nodes represent incoming edges from four different directions. In this case, at node **A**, there is a state for the dotted path coming from the north  $A_N$  and a state for the solid path coming from the west  $A_W$ . Neither of them is thrown away at Node **A** because they are distinct nodes. Both need to be relaxed and kept. Therefore, the correct shortest path, which is the dotted path, will be found. Nodes that do not have incoming edges from a particular direction will never be reached, but this doesn't pose a problem. The algorithm would ignore those nodes and never try to extract them from the heap.

Dijkstra's Routing Algorithm is modified as described above. In the worst case, there are 4 times as many nodes and edges to explore. The complexity of Dijkstra's Routing Algorithm after this modification is  $4 \times V \times \text{extract\_min} + 4 \times E \times \text{decrease\_key} = \Theta(4V \log V + 4E) = \Theta(V \log V + E)$ . The order of growth of the algorithm remains the same, still polynomial.

### 3.5 Modified Dijkstra: 1 Iteration vs. 100 Iterations

Our current Dijkstra's Routing Algorithm starts by routing the first source to destination pair. The residual capacities are then modified and the weights updated. Then it routes the second source to destination pair with the modified weights which are based on the routing result of the first source to destination pair. Next it routes for the third source to destination pair using weights modified by the first two flows and so on.

The problem with this is that based on the order chosen to route source to destination pair, the results can differ significantly. Consider the situation shown in Figure 3-3. The first pair is routed from  $S$  to  $D1$  and the path chosen is shown by the solid arrows. The second pair from  $S$  to  $D2$  is routed after the first and takes the path shown by the dotted arrows. In this case,  $S \Rightarrow A$  is congested, but  $S \Rightarrow A \Rightarrow D2$  gives the shortest path, because going around to  $S \Rightarrow C \Rightarrow B \Rightarrow D \Rightarrow D2$  is too expensive. We would like to avoid this situation because we create a heavily congested edge.




Figure 3-3: Routing result for routing  $S$  to  $D1$  first and  $S$  to  $D2$  second

If we routed the  $S$  to  $D2$  pair first, then  $S \Rightarrow A \Rightarrow D2$  would be the path chosen,

shown in Figure 3-4 by the dotted arrows. Then routing S to D1 Dijkstra's algorithm would be able to find a route that is less congested, but with the same number of hops:  $S \Rightarrow C \Rightarrow B \Rightarrow D1$ . As shown clearly in the example, routing order does matter.




Figure 3-4: Routing result for routing S to D2 first

The 100 iteration Dijkstra's Routing Algorithm is trying to mitigate the effects caused by routing order. The idea is to route each pair many times growing the bandwidth requirements gradually. This would mean when we come to last iteration the first pair is routed onto a graph where the edge weights are nearly at their final values. In this way no flow is given the luxury of having been routed onto an empty mesh which caused the problem above. The algorithm works as follows: At iteration 1, we route all source to destination pairs with only  $\frac{1}{100}$  of each pair's flow demand. At iteration 2, we first erase the route for the first source to destination pair and reroute this flow with  $\frac{2}{100}$  of its demand. During this second routing process, a small amount of all other flows are on the mesh so the edge weights are affected by them, but their contribution is small, as only  $\frac{1}{100}$  of the demand for each flow has been routed. Then we erase and reroute the second source to destination pair, the third, and so on, with  $\frac{2}{100}$  of its flow demand. On the third iteration, we reroute each flow with  $\frac{3}{100}$  of its demand, and so on. On iteration 100, we reroute with the full flow demand for all the source to destination pairs with  $\frac{99}{100}$  of the flow demand for all other flows already on

the mesh. This algorithm reduces the effect caused by routing order. For the example given above, the 100 iteration Dijkstra's Routing Algorithm is able to find the routes shown in Figure 3-4 even though  $S \Rightarrow D1$  is routed first.

This algorithm essentially repeats the previous Dijkstra's Routing Algorithm 100 times. Erasing each path requires modifying the weights of all the edges that were routed in the last iteration, and is an  $O(1)$  operation. The complexity of the algorithm is  $100 \times (V \times \text{extract\_min} + E \times \text{decrease\_key}) = \Theta(100 \times (V \log V + E)) = \Theta(V \log V + E)$ . The modification made to the 100 iteration Dijkstra's Routing Algorithm does not change the order of growth. This is still a polynomial algorithm.

### 3.6 Output Selection Process

Dijkstra's Routing Algorithm produces routes by avoiding congested links at different levels. The more congested the link is, the less likely Dijkstra's Routing Algorithm will use that link as part of a route. Given a set of values for  $C$  and  $K$ , we run Dijkstra's Routing Algorithm as described in this Chapter, taking into consideration all 20 Turn Model cases listed in Table 3.1. We reduce  $C$  from the MCL produced by XY to the MCL produced by MILP or until we cannot obtain a set of routes, storing the routes obtained for each value of  $C$ . Often times we have several routes with the same lowest MCL. We pick the set of routes with the least congestion amongst all computed routes with the lowest MCL. Congestion corresponds to the product of the average excess bandwidth demand over all links times the average number of flows competing for each link.

We developed 2 relevant metrics we can use to analyze the results.

- **Average Channel Load (ACL)**  $= \frac{\sum_{i=0}^n \text{link load for link } i}{n}$  where  $n$  is the total number of links that are used. This parameter is calculated by summing up link load for all links, and then dividing it by the number of links in the network. ACL reflects how congested the network traffic is.
- **Average Flow Competition Per VC (AFC)**<sup>2</sup>  $= \frac{\sum_{i=0}^n \text{number of flows into the VC}}{n}$  where  $n$  is the total number of virtual channels that are assigned flows to in the network. This parameter is calculated by computing all the flows going into every virtual channel, and then averaging it over all the virtual channels. AFC reflects head-of-line blocking.

One case of head-of-line blocking is illustrated in Figure 3-5. There are two flows, depicted as A and B in Figure 3-5, sharing the same link. Ahead of flow A, the traffic is heavy and that part of the network is very congested. A cannot move because the flow is congested. At the same time, there is almost no traffic ahead of the other flow B. B should be able to move freely. However, B cannot move either, because A and B

---

<sup>2</sup>VC stands for virtual channel. Each link in our mesh network should have four virtual channels, each for one direction

packets are interleaved, thus flow A is blocking flow B. The bigger the AFC, the more potential there is for head-of-line blocking.

From the above two network parameters, congestion is calculated as  $\text{congestion} = (\text{ACL} - \text{Threshold}) \times (\text{AFC})$ . **Threshold** is a parameter we can set. It can be thought of as the link capacity of the network.  $(\text{ACL} - \text{Threshold})$  gives us the average excess bandwidth demand. This number multiplied with **AFC**, which is the potential for head-of-line blocking, gives us the congestion number.

We calculate this congestion number for all generated routes with the lowest MCL, and pick the route with the lowest congestion number.




Figure 3-5: Head-of-line blocking illustration. There is heavy congestion ahead of flow A, and no congestion ahead of flow B after the split.



## Chapter 4

# Experiment and Performance: Dijkstra's Routing Algorithm

As clearly described in Chapter 3, the complexity of Dijkstra's Routing Algorithm is  $\Theta(V \log V + E)$ , which is polynomial. MILP that is introduced in [26] is an NP-hard problem, and thus an exponential algorithm. Besides, by only minimizing MCL may not give us better results than selecting routes based on each link's residual capacity, which is done by Dijkstra's Routing Algorithm. Dimension-order routing like XY and YX are simple algorithms which take linear time to route, but they may not give good routing results because the bandwidth demand of each route is not taken into account. The complexity of the three algorithms are  $XY < \text{Dijkstra's Routing Algorithm} < \text{MILP}$ .

This chapter details the simulator and simulation method we use to measure the performance in Section 4.1. Section 4.2.1 evaluates the actual performance when compared with DOR (dimension-order routing) and maximum channel load results compared with ROMM and Valiant. The analysis for the simulation results is discussed in Section 4.2.2.

## 4.1 Simulation Process

### 4.1.1 The Simulator

A cycle-accurate network simulator [12] is used to estimate throughput and latency of each flow in the application for various oblivious routing algorithms. The simulator models the router micro architecture with a 4-stage pipeline. As discussed in Section 2.3, our routing scheme only requires minor changes in modern router hardware design. Therefore, we assume an identical clock frequency and pipeline stages for all routing algorithms.

We use an  $8 \times 8$  2-D mesh network with 1 to 8 virtual channels per port. The simulator is configured to have a per-hop latency of 4 cycles, a credit latency of 1 cycle, and 2 flits of buffer size per VC. The flit length is 64 bits, and we simulate a fixed packet length of 8 flits. The buffering at each source is set to 250 packets. We expanded the number of ports between each node and each switch from 1 to 4. This way each switch can serve 4 times as many flits in each cycle. The bandwidth at each receiving node is effectively increased by a factor of 4. For each simulation, the network is warmed up for 10,000 cycles prior to being measured after which the simulation is run for 100,000 cycles to collect statistics. This is enough for convergence.

We vary the input injection rate and measure the output delivery rate. A high output delivery rate under a high input injection rate for a specific set of routes shows that the routes perform better. When the input injection rate is low, the output delivery rate should converge to the same value, dominated and controlled by the injection rate of the input. When the input injection rate is high, the network is heavily congested. A better route with a lower MCL should produce a higher output delivery rate.

Input injection rates are tuned so the maximum transmission rates of the network would be in between the two MCLs of the two routes being compared. Any difference in performance between two routes should be most visible within this region. Above this region, both networks are congested whereas below neither would be. Only within this region should the set of routes with a lower MCL not be congested while the other

set is.

### 4.1.2 Benchmarks

We use a set of standard synthetic traffic patterns, namely transpose, bit-complement, and shuffle, in our experiments. They are defined as follows:

- **bitcomp**:  $d_i = \bar{s}_i$
- **shuffle**:  $d_i = s_{i-1 \bmod b}$
- **transpose**:  $d_i = s_{i+\frac{b}{2} \bmod b}$

where  $s_i(d_i)$  denotes the  $i^{th}$  bit of the source(destination) address.

The bit length of an address is  $b = \log_2 N$ , where  $N$  is the number of nodes in the network. The synthetic benchmarks provide basic comparisons between Dijkstra's Routing Algorithm and other oblivious routing algorithms since they are widely used to evaluate routing algorithms. In the above three synthetic benchmarks, all flows have the same average bandwidth demands. We also report results on H.264, which is a widely known and used application. We profiled a 4-engine H.264 decoder and recorded all the flows' demands. This benchmark exhibits significantly varying flow demands; up to 2 orders of magnitude are possible. It is usually the case that applications have vastly different flow demands.

## 4.2 Performance Evaluation of Dijkstra's Algorithm

### 4.2.1 Channel Load Comparison

The following table 4.1 shows the MCL computed by each routing algorithm.

Traffic	XY	YX	ROMM	Valiant	Dijkstra's Routing Algorithm
transpose	175	175	200	175	75
bit-complement	100	100	400	200	100
shuffle	100	100	150	200	75
H.264	214.48	365.51	336.14	352.98	124.54

Table 4.1: Comparison of Maximum Channel Load (MCL) in MB/second.

As shown, Dijkstra's Routing Algorithm always finds routes with a lower MCL compared with XY/YX, ROMM, and Valiant routings. Dijkstra's Routing Algorithm finds the same lowest MCL as XY/YX for bit-complement. This is because the flows in bit-complement have many source and destination positions across the mesh, making it very symmetric both in the x and y direction. Figure 4-1 shows bandwidth demand for each flow and the routes determined by Dijkstra's Routing Algorithm for bit-complement. Each arrow represents a unit of flow demand. The more arrows pointing to the same direction on a link, the more congested the link is on that direction. Symmetry is presented clearly on the 2-D mesh. For comparison, the same plot is generated for Transpose, shown in Figure 4-2. There is not as much symmetry in this figure. For the case of bit-complement, this symmetry makes a minimal and deterministic routing algorithm, XY/YX, attractive for this particular traffic pattern.

However, Dijkstra's Routing Algorithm finds the same result as XY/YX for bit-complement, which shows its advantage. Dijkstra's Routing Algorithm produces routes with the same or lower MCL than XY/YX.

ROMM and Valiant produce routes that have higher MCLs than XY/YX for the 3 synthetic benchmarks. The 3 synthetic benchmarks that we use have relatively symmetric traffic patterns, which makes XY/YX attractive. Even though ROMM and Valiant add randomness into the network to balance the load, the MCLs of the routes ROMM and Valiant produce are higher than routes produced by XY/YX for

bit-complement, transpose and shuffle. It should be noted that this is not the case for H.264.




Figure 4-1: Routes generated using Dijkstra's Routing Algorithm for Bit-complement.

#### 4.2.2 Simulation Results of 4 Synthetic Benchmarks

Routes are produced by Dijkstra's Routing Algorithm on 4 benchmarks, discussed in Section 4.1.2. Since there are 20 different Turn Models described in Section 3.3, there are 20 results produced. The selection process described in Section 3.6 is used to pick the routes with the lowest MCL and least congestion. It selects one set




Figure 4-2: Routes generated using Dijkstra's Routing Algorithm for Transpose.

of routes for each benchmark. The threshold is set to be 25 for all 4 benchmarks. Figures 4-3 through 4-6<sup>1</sup> show the output delivery rate corresponding to various input injection rates for Dijkstra's Routing Algorithm, XY, and YX on the 4 benchmarks. Normalized injection rate is the injection rate per MB/s bandwidth.

As can be observed from the 4 Figures 4-3 through 4-6, Dijkstra's Routing Algorithm out-performs existing oblivious routing algorithms significantly in most of the benchmarks. This is seen in the reduction of the MCL and the number of congested links whose channel load is close to the MCL. Figures 4-4, 4-5, and 4-6 clearly show when the injection rate is low, the output delivery rate is dominated by the injection rate. The delivery rate converges to the same value for all routing algorithms.

An interesting result to note is that of bit-complement. The best result is given by minimal routing, YX. This is because the flow demands in bit-complement are very symmetric, as described in Section 4.2.1. This is the reason why YX routing would give us the lowest MCL, and thus a good simulation result. This is not normally the case. It happened to be true for bit-complement. It is verified that Dijkstra's Routing Algorithm can produce the same routes as XY or YX when XY/YX routes are optimal.

### 4.2.3 Simulation Result with Bandwidth Variations

The motivation for this experiment stems from the fact that the bandwidth profiling is unlikely to be exact, it is important to know whether the routes produced by Dijkstra's Routing Algorithm are still robust enough to out-perform DOR. Figure 4-7 shows the performance when after the bandwidths have been varied. In this experiment, we use the same routes as those used in Figure 4-4. The demands are, however, changed to be between  $\pm 10\%$  to  $\pm 50\%$  of the original. For each flow, the bandwidth could be as little as half or as much as 1.5 times as the original demand. It is clear that Dijkstra's Routing Algorithm continues to out-perform XY/YX because it spreads the load to achieve a global optimal set of routes.

---

<sup>1</sup>XY-1 is the performance of XY with 1 virtual channel, YX-2 is YX with 2 virtual channels and DJK-2 is Dijkstra's Routing Algorithm with 2 virtual channels, etc.

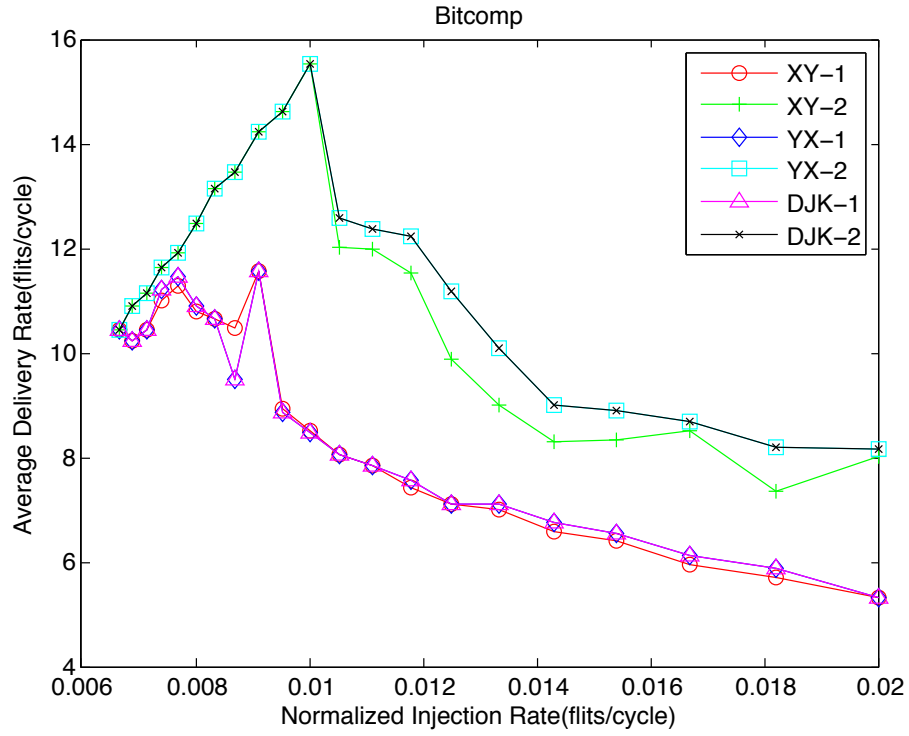


Figure 4-3: Load-throughput graphs for bit-complement on a router with 1 or 2 VCs.




Figure 4-4: Load-throughput graphs for transpose on a router with 1 or 2 VCs.



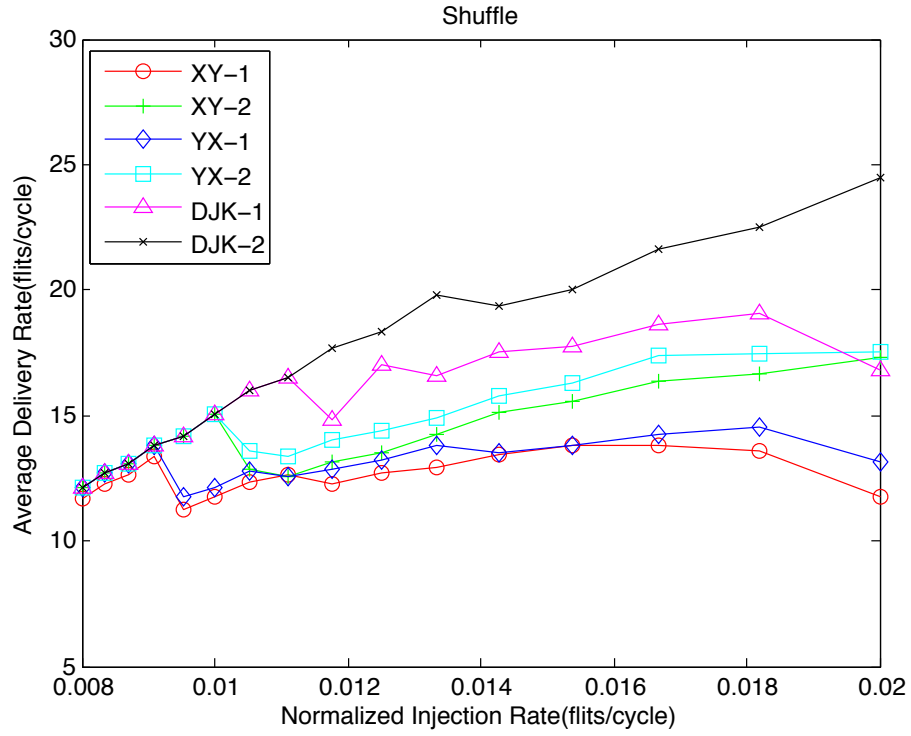


Figure 4-5: Load-throughput graphs for shuffle on a router with 1 or 2 VCs

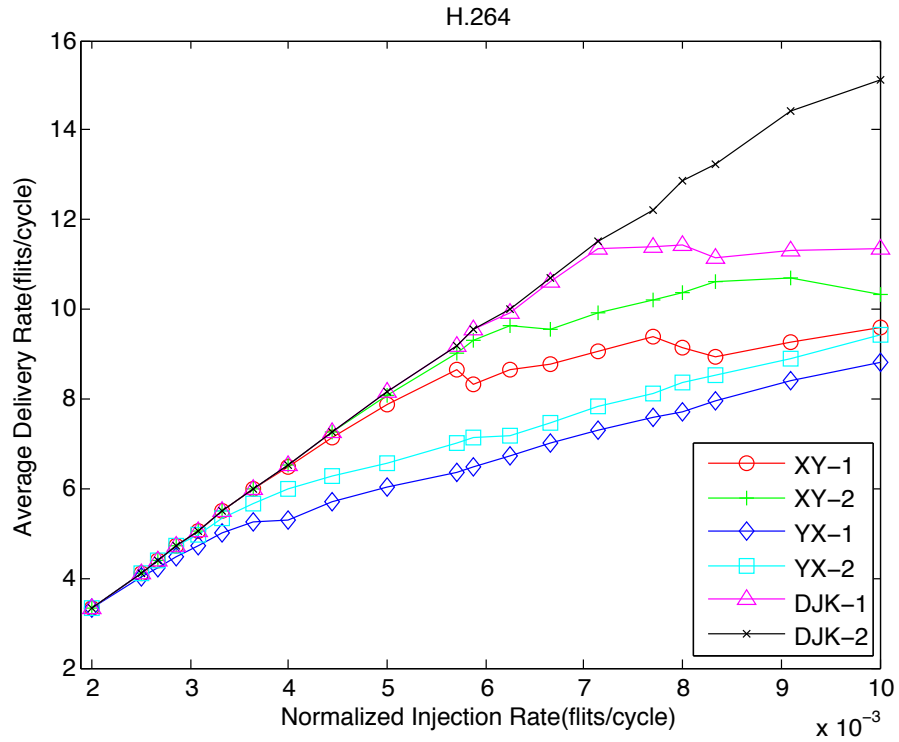


Figure 4-6: Load-throughput graphs for H.264 on a router with 1 or 2 VCs.

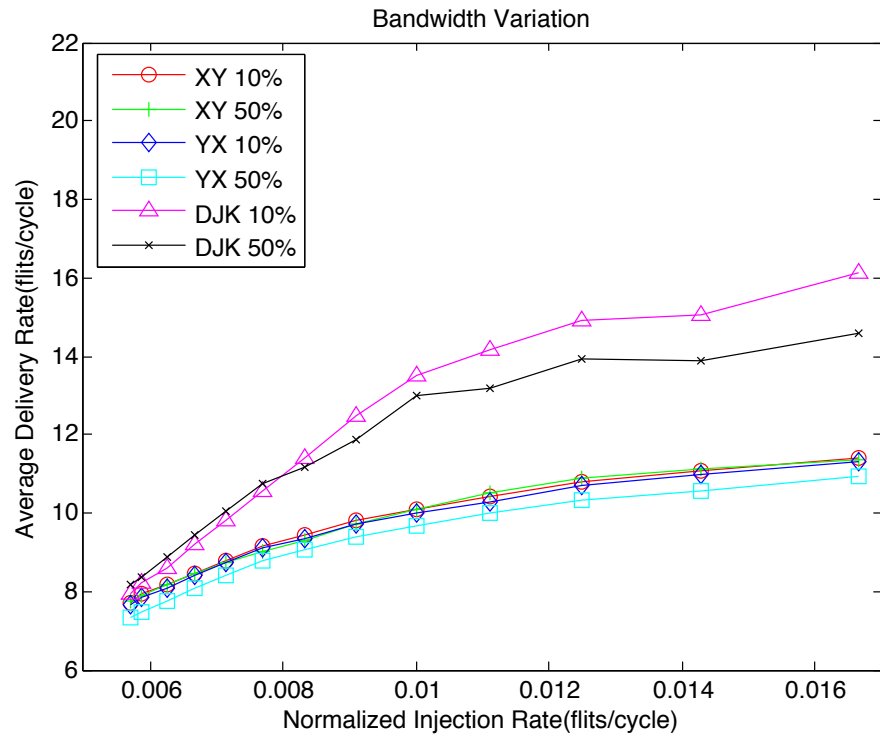


Figure 4-7: Load-throughput graphs for transpose (1 virtual channel) when bandwidths change by  $\pm 10\%$  and  $\pm 50\%$  after route computation.

## Chapter 5

# Dijkstra's Routing Algorithm in Minimal Routing

After generating a set of routes, flows need to be allocated to virtual channels to reduce sharing among them and improve performance. It turns out that VC allocation can help in avoiding deadlock as well if the routes satisfy some conditions. Section 5.1 discusses motivations to modify Dijkstra's Routing Algorithm to do minimal routing. Section 5.2 details the VC allocation procedures following the routing step. Section 5.3 talks about detailed modifications needed from the original Dijkstra's Routing Algorithm introduced in Chapter 3 in order to generate bandwidth-sensitive minimal routes. Section 5.4 details the method Dijkstra's Routing Algorithm in Minimal Routing uses to favor certain minimal routes. Section 5.5 summarizes the four combinations of turns we can favor in order to make the VC allocation step more efficient.

## 5.1 Motivations for exploring Dijkstra’s Algorithm in Minimal Routing

Dijkstra’s Routing Algorithm described in Chapter 3 guarantees us a set of deadlock free routes, which have a low MCL. The head-of-line blocking problem introduced in Section 3.6, though is taken into account during the output selection process, is not considered heavily during the routing process. After the routing process, we need to do virtual channel allocation based on the number available, which is constrained by the hardware. We allocate virtual channels in the way that reduces head-of-line blocking. This is described in detail in Section 5.2.

Virtual channel allocation can not only reduce head-of-line blocking problems, but also break some cyclic routes ensuring they are deadlock free. ROMM and Valiant routings, for example, require at least 2 virtual channels to avoid deadlock. The first segment of each route before the intermediate point needs to be assigned to one virtual channel, and the second segment is assigned to the other. In this way, routes produced by ROMM and Valiant do not cause deadlock problems. If only one virtual channel is used, ROMM and Valiant routes do not guarantee a deadlock free environment.

With the aid of VC allocation, Dijkstra’s Routing Algorithm need not generate deadlock free routes. As long as deadlock problems can be solved with a limited number of virtual channels, we can ensure the routes can be deadlock free after VC allocation.

As described in [23], we only need 2 sets of virtual channels to ensure deadlock freedom if all routes are minimal. If there are more than 2 virtual channels available, some of them can be dedicated to reducing head-of-line blocking. This motivates us to explore generating bandwidth-sensitive minimal routes using Dijkstra’s Algorithm.

## 5.2 Virtual Channel Allocation for Dijkstra's Algorithm in Minimal Routing

Dijkstra's Algorithm in Minimal Routing produces routes that are minimal but not deadlock free. At least 2 virtual channels need to be dedicated to make the routing deadlock free. If we have more than 2 VCs, we divide all the virtual channels into virtual channel sets. 1 VC set has to be dedicated to the West-First Turn Model while the other has to be dedicated to the North-Last Turn Model. All the routes produced by Dijkstra's Algorithm in Minimal Routing are minimal, and therefore conform to one of the 2 Turn Models, West-First or North-Last, or both. We group all the routes into 3 sets. The first set is for routes that only fit into the West-First Turn Model, the second for routes that only fit into the North-Last Turn Model while the third set is for routes that fit into both Turn Models. Therefore, at each link, we may have non-zero routes in the first and second set. In this case, any of the routes in the third set should be distributed into either of the first two sets, carrying the strategy of reducing head-of-line blocking and balancing the load, i.e.: the number of flows per VC. If either the first or the second set is empty or both are empty, then there is greater flexibility to reduce head-of-line blocking and balance number of flows per VC. In other words, the more routes the third set has, the more choice we have in assigning VCs at each link, because every flow in the third set can go into either of the two VCs.

VC allocation is carried out to best reduce head-of-line blocking and to balance the number of flows per VC. The details of the procedure are described in [23]. The following is a summary of the steps.

1. The algorithm searches for a VC that only contains flows which have shared a VC with the flow that is currently being allocated a VC. If this VC exists, the flow is assigned to the VC. This check is done because there is no advantage in assigning another VC to this new flow since head-of-line blocking may already exist between the new flow and all the flows that were in the VC because they have shared a VC before.

2. The algorithm looks for an empty VC. If found, the new flow is assigned to this empty VC.
3. The algorithm looks for a VC that contains a flow that has shared a VC with the flow that is to be assigned a VC. If the VC is found, the new flow is assigned to the VC.
4. The algorithm assigns the new flow to the VC that contains the minimum number of flows.

Sharing between flows for this particular VC is updated after each iteration. This procedure above is repeated for each flow utilizing the given link. And this is repeated for each link.

## 5.3 Implementation of Bandwidth-Sensitive Minimal Routing using Dijkstra's Algorithm

There are a few implementation changes which need to be made to Dijkstra's Routing Algorithm described in Chapter 3 in order to generate minimal routes.

### 5.3.1 Elimination of Turn Model Constraint

First of all, Turn Model does not need to be used because our final routing does not need to be deadlock free. Shim's paper [23] proves that a set of minimal routes can be made deadlock free by using 2 virtual channels. There are 8 different combinations of turns that can appear in minimal routes given a set of source to destination pairs on a 2-D mesh. We can assign 4 of them to one virtual channel, which contains all the routes that satisfy the West-First Turn Model, shown in Figure 5-1. We then assign the other 4 of them to the other virtual channel, which contains all the routes that satisfy the North-Last Turn Model, shown in Figure 5-2. Therefore, we also do not need to keep track of incoming and outgoing directions during execution.




Figure 5-1: The West-First Turn Model.

### 5.3.2 Elimination of Four States Per Node

Since we do not need to keep track of the directions at each node along a particular route, we do not need to keep four states per node. The reason why these were needed in the previous design was because at each node, we wanted to keep a separate state




Figure 5-2: The North-Last Turn Model.

for the routes from all four directions, even though some of them had a larger weight up to that point the current shortest path may be eliminated by the Turn Model later on, due to certain prohibited turns on the path after the node. This is illustrated in Section 3.4.

### 5.3.3 The Addition of Making Some Progress at Every Step

We want our routes to be minimal, which means the route has the smallest number of hops possible. In order to guarantee minimal routing, we enforce checks of the distance between the current node and destination frequently. At each node, we check the distance left to the destination, and designate this  $d1$ . For every node adjacent to the current one, we calculate the distance between this node and the destination and call this  $d2$ . If  $d1 > d2$ , the node is a potential candidate for our minimal route and should be relaxed if needed. If not, the partial route cannot make a minimal route and thus invalid. In this case, we add `positive_infinity` to the weight of the edge to make sure it is never relaxed or chosen by our routing algorithm.

### 5.3.4 The Need of 100 Iterations

Though constrained to minimal routing, there are still many viable paths from a source to a destination. Figure 5-3 shows 4 possible minimal routes from S to D where S and D are 2 units in length and 2 units in width apart from each other. Given many choices of routes for each source to destination pair, the route chosen



will affect the routing process of the next flow. The route that is then chosen for this flow will affect the routing process of the next, and so on. In other words, the routing order still matters in this case. This is the reason why 100 iterations are still desired and needed to reduce the effect of the routing order and to achieve a set of globally optimal routes.




Figure 5-3: 4 possible minimal routes from S to D where S and D are 2 units in length and 2 units in width apart from each other.

## 5.4 Dijkstra's Routing Algorithm in Minimal Routing with Favored Turns

Shim's paper [23] discusses that if we are using West-First and North-Last Turn Models as shown in Figure 5-1 and 5-2, there are 4 minimal routes that fit into both Turn Models. These 4 turns are shown in Figure 5-4. The idea is to get as many minimal routes that are solely composed of these four turns as possible as these minimal routes can be assigned to either virtual channel. (Given two virtual channels, one virtual channel has all the routes satisfying the West-First Turn Model, and the other has routes satisfying North-Last Turn Model.) The more of these minimal routes we have, the more routes we have where we can choose which VC to assign it to. And so we can dedicate more effort to reducing head-of-line blocking by reducing the amount of sharing and splitting.




Figure 5-4: The four (out of possible eight) different one-turn routes on a 2-dimensional mesh that conform to both the West-First and North-Last Turn Model.

We call the routes that are solely composed of any of the four turns shown in Figure 5-4 favored routes. Note that the favored routes can only make one turn. In Figure 5-5, the solid path is a favored route because it only contains the east-north turn. The dotted path, however, is not a favored route because it also makes a north-east turn. Favored routes (solid path) can only make one turn along the path from

source to destination. A route that zig-zags (dotted path) can never be a favored route.




Figure 5-5: Two routes to illustrate favored routes. The solid route is a favored route while the dotted route is not.

#### 5.4.1 Implementation For Favoring Four Turns

We want to end up with a set of routes that have as many favored routes as possible, but at the same time a low MCL and low congestion. We do not wish to modify the weight of a favored route. If we did, the weights would be determined both by the residual capacity of the link and the position of the nodes the route has gone through, two completely unrelated criteria. How the weights should be reduced in order to give priority to the favored route is not clear. We do not want to compromise MCL and low congestion to get more favored routes. Instead, we only want to give priority to those favored routes. Given a source node and a destination node that are not in the same row or column, we can only find one possible favored route. We calculate and update the residual capacities and weights along all links traversed by this favored route. We then relax and put each node into the heap in the same way as Dijkstra's Routing Algorithm.

With all the nodes and edges along this path already explored, we run Dijkstra's Algorithm in Minimal Routing as we did before: exploring all possible minimal routes from source to destination; if the condition for relaxation is satisfied, that is, if the sum of the weight of the parent node and the weight of the edge is smaller than the current weight of the node, then relax that edge by setting the weight of the current node to be the new smaller weight and change the parent node pointer. Finally change the weight of the node in the heap. We keep doing this for all possible minimal paths. This method gives priority to favored routes. In the case of having two minimal paths with same lowest weight, if one of which is a favored route, then this favored route is guaranteed to be chosen by this implementation. The reason is that all the distances of the nodes on this favored route have been initialized. Later on if there is another path that results in the same lowest weight at the destination, the destination node will not be relaxed and its parent pointer will still use the favored route.

The favored route is only chosen when it does not contribute any more to congestion given the flows that are already routed on the mesh. A favored route will not be chosen if there is another path that can result in a lower weight at the destination. In this case the path will be relaxed along the new path and our algorithm will output the path with the lowest congestion.

### 5.4.2 Advantage of this Implementation

The initialization step saves us from having to constantly relax and update the heap. If we did not initialize the graph with the favored minimal route, and instead we ran Dijkstra's Algorithm in Minimal Routing without this initialization many more checks would be required. Every time an edge is to be relaxed the algorithm would proceed normally, updating the weight parent node pointer appropriately. The extra check is needed when the sum of the weight of the parent node and the weight of the edge is the same as the weight of the current node. Here we must check if the new route is a favored route and if so, we need to relax the edge to replace the old route with this new favored route. There is no need to update the weight as it stays the same. This implementation is not as efficient, because we have to relax every time we

see a smaller weight at a node less than the current weight of the node. Furthermore, every time we find the weights are equal and we have a choice over routes we would have to check the position of every node on the route to decide if the new route is a favored one. With initialization, if the preferred path has the lowest weight, we do not need to relax at all. Even if the preferred path does not have the lowest weight, we may still save ourselves from relaxing and updating as frequently.

## 5.5 Four Turn Model Pairs for Minimal Routing

The West-First and North-Last Turn Model pairs can be rotated four times, each time with different turns being disallowed. Figure 5-6 shows the four Turn Model pairs and the favored minimal routes that satisfy both the West-First and North-Last Turn Model.

	WF	NL	Favored Routes
no rotation			
rotate counterclockwise once			
rotate counterclockwise twice			
rotate counterclockwise 3 times			

Figure 5-6: West-First and North-Last Turn Models rotated four times to form 4 Turn Model pairs. The column WF shows the West-First Turn Model rotated 4 times. The column NL shows the North-Last Turn Model rotated 4 times. The last column shows the favored minimal routes that can be assigned to either virtual channel.

Since we can look at the Turn Model pairs four different ways and the set of favored turns are different in every case, we should run Dijkstra's Algorithm in Minimal Routing four times based on the four orientations of the Turn Model Pairs. For each choice of Turn Model pair, we should have more favored routes where we have a choice in VC allocation than if we do not favor any routes.

## 5.6 Routes Selection for Dijkstra's Algorithm in Minimal Routing

We should also run Dijkstra's Algorithm in Minimal Routing with no favored turns applied. This way we end up with five sets of routes: four sets of routes with favored turns, and one set of routes with no favored turns. The selection process is done the same way as described in Section 3.6. We pick the routes with the minimum MCL and the lowest congestion.





## Chapter 6

# Performance Comparison Between Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing

Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing are two ways to solve a routing problem. Dijkstra's Routing Algorithm tries to reduce congestion of the network and produces a set of deadlock free routes. Dijkstra's Algorithm in Minimal Routing also tries to reduce congestion of the network, but produces minimal routes that are not deadlock free. The routes need to go through virtual channel allocation, which reduces head-of-line blocking, and requires at least 2 VCs to make the simulation deadlock free. It is interesting to see how these two routing algorithms perform against each other. Simulation results are shown in Section 6.2 for the four benchmarks when 2 VCs are used in both routing algorithms. Performance is also compared and analyzed when more VCs are available. Simulation results are shown in Section 6.3 for Shuffle and H.264 with multiple VCs for both routing algorithms.

## 6.1 Channel Load Comparison

The MCL for routes produced by Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing, is shown in Table 6.1.

Traffic	Dijkstra's Routing Algorithm	Dijkstra's Algorithm in Minimal Routing
transpose	75	75
bit-complement	100	125
shuffle	75	75
H.264	124.54	174.07

Table 6.1: Comparison of MCL in MB/second for two variations of routing algorithms based on Dijkstra's Algorithm.

## 6.2 Performance Comparison of 2 VCs on Four Benchmarks

With the simulator configured the same way as described in Chapter 4, simulations were run for Dijkstra's Routing Algorithm with dynamic VC allocation and Dijkstra's Algorithm in Minimal Routing with static VC allocation with 2 virtual channels being used for the four benchmarks introduced in Section 4.1.2. We chose 2 virtual channels as this is the smallest number that allows Dijkstra's Algorithm in Minimal Routing to produce deadlock free routings. The static virtual channel allocation process described in Section 5.2 with 2 VCs is used for the routes produced by Dijkstra's Algorithm in Minimal Routing to make the simulation deadlock free. The performance is shown in Figures 6-1 to 6-4.<sup>1</sup>

As can be observed, the performance of Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing are very close to each other. They both significantly out-perform XY/YX. This is not surprising, because they both iteratively and heuristically reduce heavily congested links.

---

<sup>1</sup>DJK MIN is the simulation result for Dijkstra's Algorithm in Minimal Routing. DJK is the simulation result for Dijkstra's Routing Algorithm.




Figure 6-1: Performance comparison for bit-complement on a router with 2 VCs.

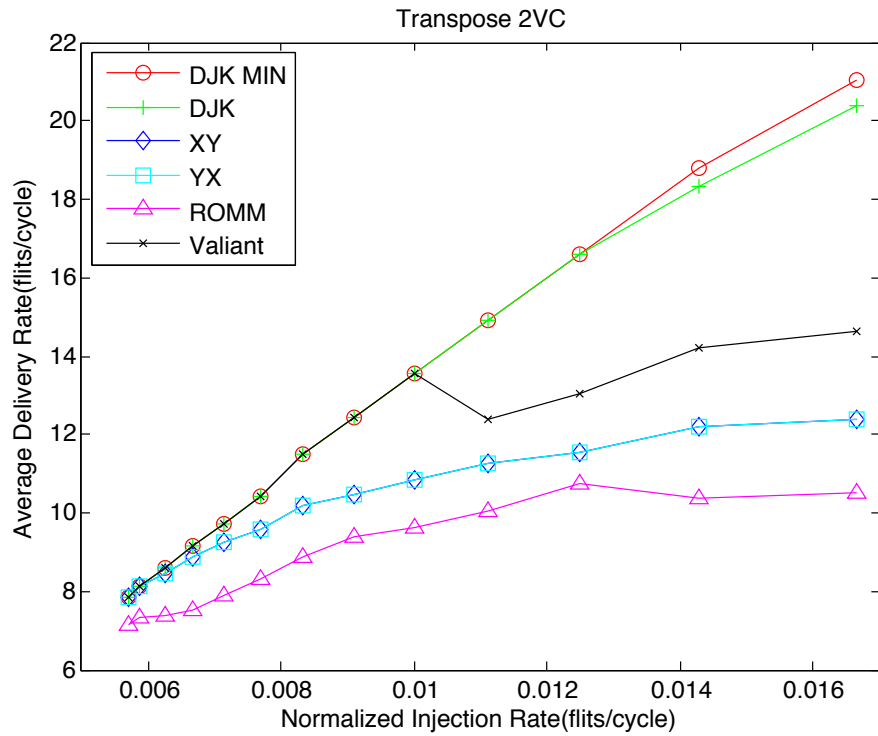


Figure 6-2: Performance comparison for transpose on a router with 2 VCs.

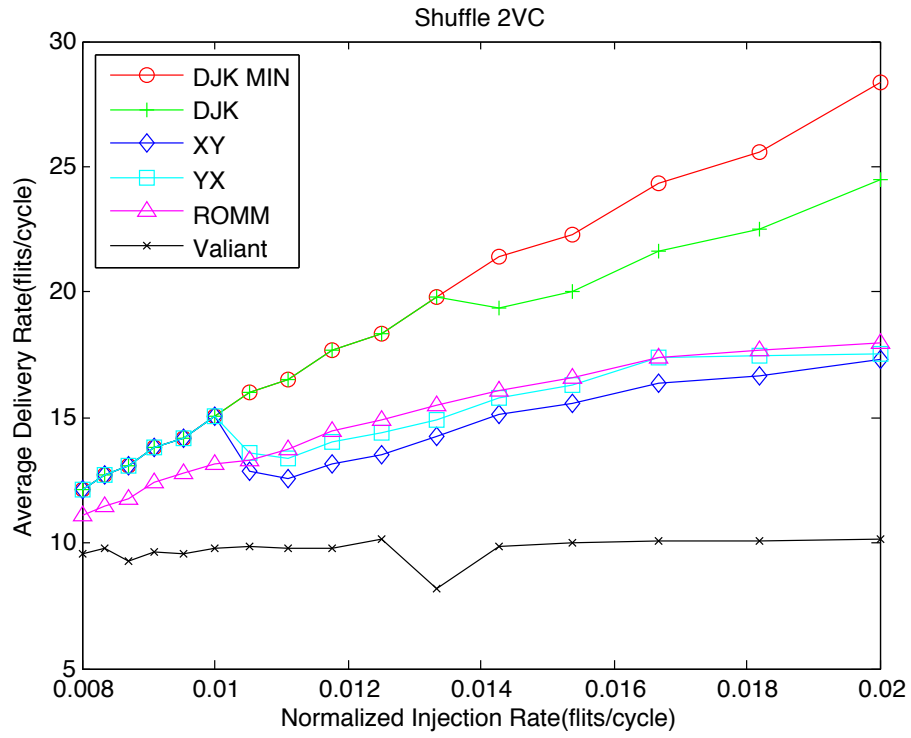


Figure 6-3: Performance comparison for shuffle on a router with 2 VCs.

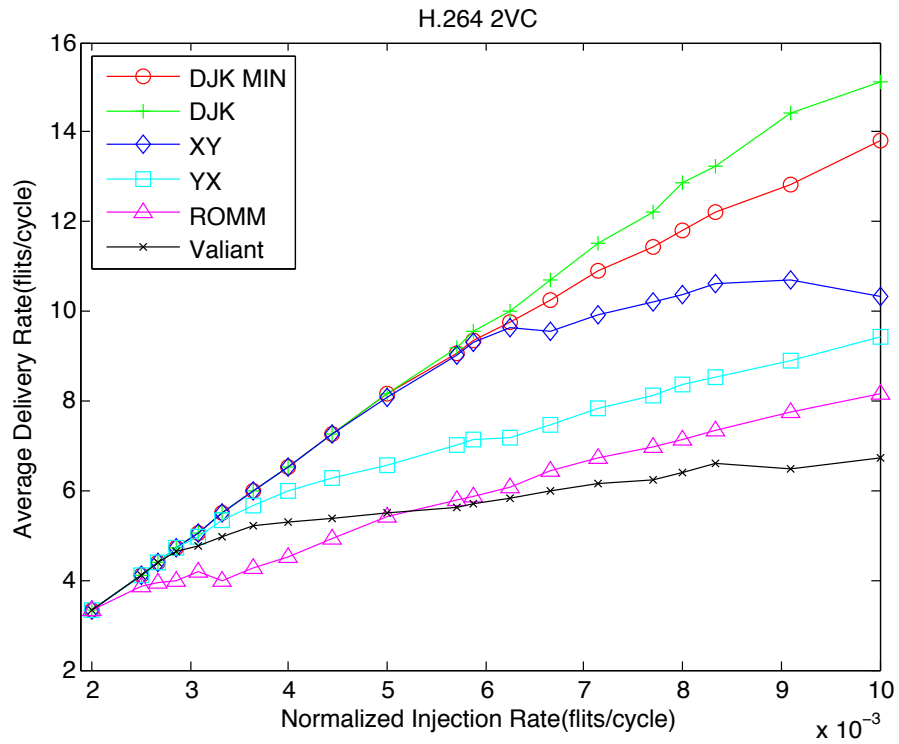


Figure 6-4: Performance comparison for H.264 on a router with 2 VCs.

It can be observed that static VC allocation on the routes produced by Dijkstra's Algorithm in Minimal Routing performs better than the dynamic VC allocation used on routes produced by Dijkstra's Routing Algorithm. This is because static VC allocation reduces head-of-line blocking much more effectively than dynamic VC allocation. In the case shown in Figure 6-2 and 6-3 when both the routes produced by Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing have the same lowest MCLs, namely, 75, Dijkstra's Algorithm in Minimal Routing with static VC allocation with 2 VCs performs better.

For the case of Bit-complement, as shown in Figure 6-1, Dijkstra's Algorithm in Minimal Routing produces a set of routes with a MCL of 125, which is larger than the MCL of the routes produced by Dijkstra's Routing Algorithm. Therefore, when the network is not congested, Dijkstra's Routing Algorithm performs better because there is not much head-of-line blocking when there is light traffic, so the performance is dominated by MCL (lower MCL corresponds to better performance); On the other hand, when the network is heavily congested, Dijkstra's Algorithm in Minimal Routing with static VC allocation performs better. Since static VC allocation reduces head-of-line blocking more effectively, Dijkstra's Algorithm in Minimal Routing performs better under congested network traffic.

For the case of H.264, as shown in Figure 6-4, Dijkstra's Routing Algorithm, which has an MCL of 124.54, consistently performs better than Dijkstra's Algorithm in Minimal Routing, which has an MCL of 174.07. In this case, the difference in head-of-line blocking through VC allocation for the two routes becomes minor, and the significant difference between the two MCLs becomes the dominating factor. Therefore, Dijkstra's Routing Algorithm performs better for every injection rate for H.264.

One thing to note is that all routes produced by Dijkstra's Routing Algorithm for the three synthetic benchmarks (Figure 6-1, 6-2, 6-3) are minimal even without the minimal constraint during routing. The fact that routes produced by the two algorithms being similar corresponds well with similar simulation results.

## 6.3 Performance Comparison of multiple VCs

### 6.3.1 Shuffle with 2 and 4 VCs

Performance is compared between Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing with static VC allocation on the Shuffle benchmark with virtual channel numbers being 2 and 4. The simulation results are shown in Figure 6-5.

Note that static VC allocation is applied to routes produced by Dijkstra's Routing Algorithm. Since the routes are deadlock free by themselves, we only need to apply the VC allocation method described in Section 5.2 to reduce head-of-line blocking.

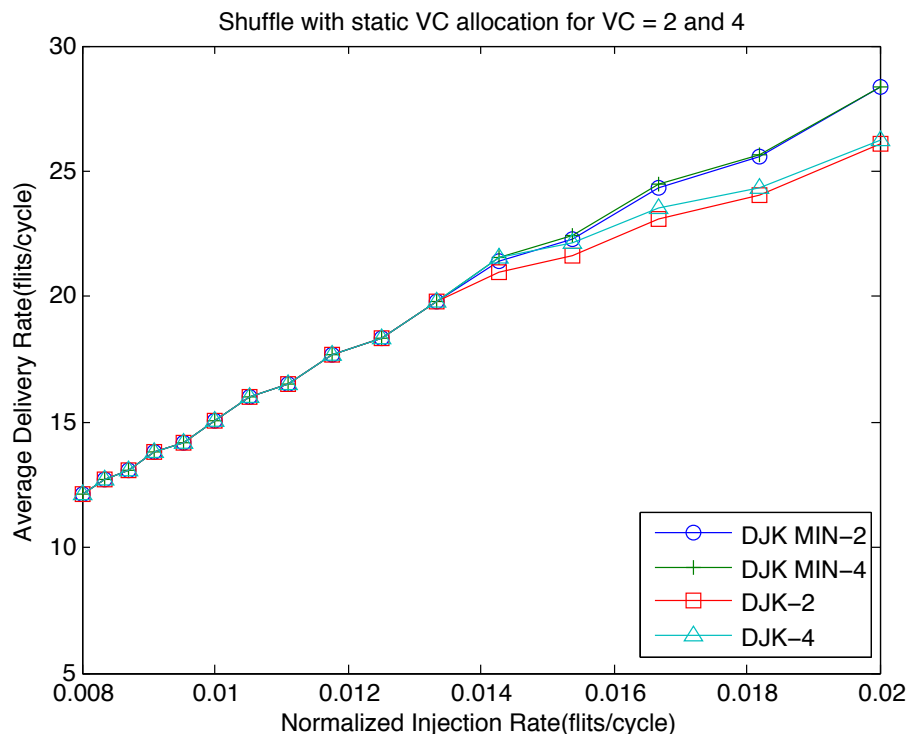


Figure 6-5: Performance comparison for Shuffle on a router with 2 and 4 VCs.

It can be observed that Dijkstra's Algorithm in Minimal Routing performs better than Dijkstra's Routing Algorithm for  $VC = 2$ , because head-of-line blocking is slightly more severe in the routes produced by Dijkstra's Routing Algorithm. However, the performance of both algorithms for  $VC = 4$  seems to converge. When the number of virtual channels is greater or equal to the largest number of flows shared

per link, which is 3 in both cases, routes produced by either routing algorithm can have private channels for each flow. For the case of VC being 4, each flow in both algorithms is able to be allocated a private channel, and thus head-of-line blocking is completely eliminated. Therefore, both routes produced by the two algorithms after VC allocation are performing with no head-of-line blocking effects whatsoever and converge to similar results because both routes have the same MCL, namely, 75. There is no benefit in adding more virtual channels after head-of-line blocking is completely eliminated.

### **6.3.2 H.264 with private channels**

Figure 6-6 shows the simulation result with 10 VCs on XY, YX, Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing. In this experiment, all the flows have private channels, so there is no head-of-line blocking problem for any route. As can be observed, routes produced by Dijkstra's Routing Algorithm with the lowest MCL 124.54 perform the best. Routes produced by YX with the highest MCL 365.51 perform the worst. There is a direct correlation between MCL and performance when each flows has a private channel; the lower the MCL, the better the performance.

We can conclude from the above simulation result that when the number of virtual channels is too few to allow private channels, performance is controlled by MCL and head-of-line blocking. When the number of virtual channels is big enough to give every flow private channels, head-of-line blocking is completely eliminated and performance is solely dominated by MCL.

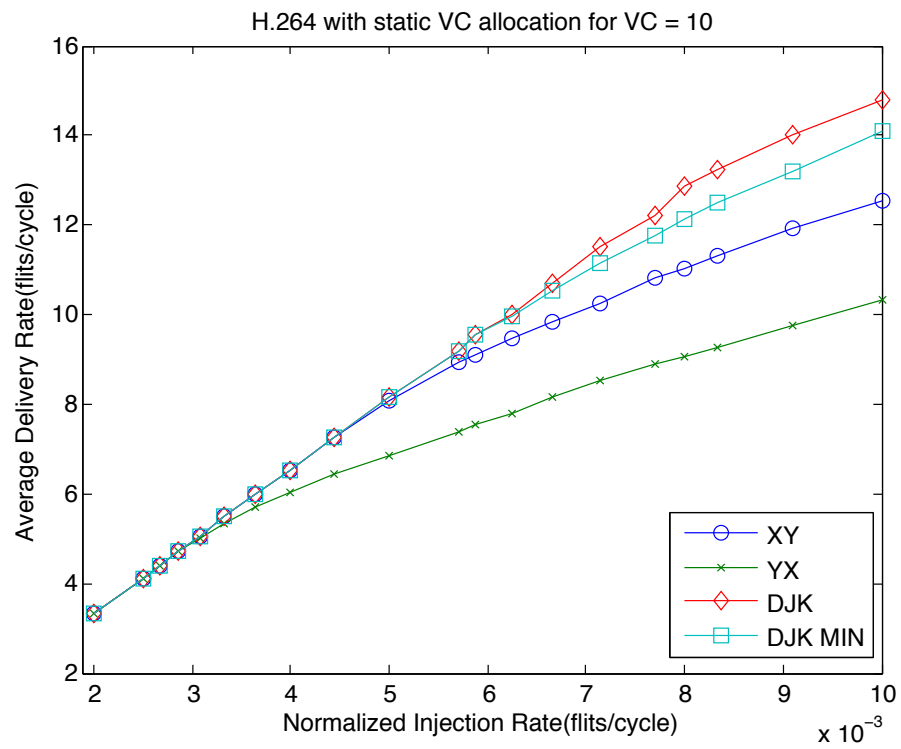


Figure 6-6: Performance comparison for H.264 on a router with 10 VCs.



# Chapter 7

## Conclusion and Future Work

Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing were experimentally shown to out-perform dimension order routings such as XY and YX. The performance advantage is obvious, but we need to acquire the knowledge of rough bandwidth requirements of the application ahead of time and also need polynomial time to compute the routes. For applications that we can profile or have the knowledge of communication requirements ahead of time, bandwidth-sensitive oblivious routing would be suitable for its obvious performance benefit.

Sometimes it is difficult to get rough bandwidth requirements for each flow. Bandwidth-sensitive oblivious routing would still perform well if we know all the high flow demand or latency critical paths, because we can just route using these paths.

Even though the performance between Dijkstra's Routing Algorithm and Dijkstra's Algorithm in Minimal Routing are similar, Dijkstra's Routing Algorithm should be used when there are less than 2 VCs or when the set of routes generated has a lower MCL than that of Dijkstra's Algorithm in Minimal Routing. Dijkstra's Algorithm in Minimal Routing should be chosen otherwise to get the best performance. In either case, static VC allocation should be used to achieve better performance by reducing head-of-line blocking.

When each flow has a private channel, the MCL is the only factor affecting performance. If we have a large number of virtual channels available to guarantee private

channels, we should concentrate all our efforts on reducing the MCL. Dijkstra's Routing Algorithm is the best method at effectively reducing the MCL.

All of the experiments in this thesis were performed using a cycle accurate simulator which only simulates simple packet generation, consumption and communications. In real applications, data has dependencies which may affect communication drastically. Future work should be done running real applications on a simulator with the ability to model data dependencies to verify the advantage of bandwidth-sensitive oblivious routing.

In addition, routing is not the only step that has to be completed before running applications. The placement of the modules would affect the routing result. Having an adaptive network may make our bandwidth-sensitive oblivious routes perform even better, or worse. Future work should be done in finding a graceful way to combine placement, routing, and building an adaptive network.

# Appendix A

## Acronyms

Term	Meaning	Defined In
ACL	Average Channel Load	3.6
AFC	Average FLow Competition Per VC	3.6
CDG	Cycle Dependency Graph	3.3
DOR	Dimension Order Routing	2.2.1
MCL	Maximum Channel Load	3.2
MILP	Mixed Integer-Linear Programming	3.2
PE	Processing Element	2.1
VC	Virtual Channel	3.6
WOT	Weighted Ordered Toggle	2.2.2



# Bibliography

- [1] T. Bjerregaard and S. Mahadevan. A survey of research and practices of Network-on-chip. *ACM Computing Surveys (CSUR)*, 38(1), 2006.
- [2] M.H. Cho. *Diastolic Arrays: Throughput-Driven Reconfigurable Computing*. PhD thesis, Massachusetts Institute of Technology, 2008.
- [3] W.J. Dally. *Principles and Practices of Interconnection Networks*. Morgan Kaufmann, 2004.
- [4] W.J. DALLY and C.L. SEITZ. Deadlock-Free Message Routing in Multi-processor Interconnection Networks. *IEEE Trans. Computers*, 36(5):547–533, 1987.
- [5] W.J. Dally and B. Towles. *Principles and Practices of Interconnection Networks*. Morgan Kaufmann, 2003.
- [6] E.W. Dijkstra. A note on two problems in connexion with graphs. *Numerische Mathematik*, 1(1):269–271, 1959.
- [7] J. Duato. A new theory of deadlock-free adaptive routing in wormhole networks. *IEEE Transactions on Parallel and Distributed Systems*, 4(12):1320–1331, 1993.
- [8] J. Duato. A Necessary and Sufficient Condition for Deadlock-Free Adaptive Routing in Wormhole Networks. *IEEE TRANSACTIONS ON PARALLEL AND DISTRIBUTED SYSTEMS*, page 10551067, 1995.
- [9] R. Gindin, I. Cidon, and I. Keidar. NoC-Based FPGA: Architecture and Routing. *NOCS07*, 2007.
- [10] C.J. Glass and L.M. Ni. The turn model for adaptive routing. *Journal of the ACM (JACM)*, 41(5):874–902, 1994.
- [11] Z. Guz, E. Bolotin, I. Cidon, R. Ginosar, and A. Kolodny. Efficient link capacity and QoS design for network-on-chip. In *Proceedings of the conference on Design, automation and test in Europe: Proceedings*, pages 9–14. European Design and Automation Association 3001 Leuven, Belgium, Belgium, 2006.
- [12] H. Hossain, M. Ahmed, A. Al-Nayeem, T. Islam, and M. Akbar. Gpnocsim - A General Purpose Simulator for Network-On-Chip. *Information and Communication Technology, 2007. ICICT '07*, pages 254–257, March, 2007.

- [13] J. Hu and R. Marculescu. Exploiting the routing flexibility for energy/performance aware mapping of regular NoC architectures. In *Design, Automation and Test in Europe Conference and Exhibition*, pages 688–693, 2003.
- [14] Michel Kinsy, Myong Hyon Cho, Tina Wen, Edward Suh, Marten van Dijk, and Srinivas Devadas. Application-Aware Deadlock-Free Oblivious Routing. *The 36th International Symposium on Computer Architecture (ISCA 2009)*, 42, 2009.
- [15] J.M. Kleinberg. *Approximation algorithms for disjoint paths problems*. PhD thesis, Massachusetts Institute of Technology, 1996.
- [16] R. Mullins, A. West, and S. Moore. Low-Latency Virtual-Channel Routers for On-Chip Networks. In *Proceedings of the 31st annual international symposium on Computer architecture*. IEEE Computer Society Washington, DC, USA, 2004.
- [17] S. Murali and G. De Micheli. SUNMAP: a tool for automatic topology selection and generation for NoCs. In *Proceedings of the 41st annual conference on Design automation*, pages 914–919. ACM New York, NY, USA, 2004.
- [18] T. Nesson and S.L. Johnsson. ROMM routing on mesh and torus networks. In *Proceedings of the seventh annual ACM symposium on Parallel algorithms and architectures*. ACN New York, NY, USA, 1995.
- [19] L.M. Ni and P.K. McKinley. A Survey of Wormhole Routing Techniques in Direct Networks. *IEEE Computer*, 26(2):62–76, 1993.
- [20] M. Palesi, G. Longo, S. Signorino, R. Holsmark, S. Kumar, and V. Catania. Design of bandwidth aware and congestion avoiding efficient routing algorithms for networks-on-chip platforms. In *Proc. of the ACM/IEEE International Symposium on Networks-on-Chip (NOCS)*, pages 97–106, 2008.
- [21] L.S. Peh and WJ Dally. A delay model and speculative architecture for pipelined routers. In *High-Performance Computer Architecture, 2001. HPCA. The Seventh International Symposium on*, pages 255–266, 2001.
- [22] D. Seo, A. Ali, W.T. Lim, N. Rafique, and M. Thottethodi. Near-Optimal Worst-Case Throughput Routing for Two-Dimensional Mesh Networks. In *International Symposium on Computer Architecture: Proceedings of the 32 nd annual international symposium on Computer Architecture*, volume 4, pages 432–443, 2005.
- [23] Keun Sup Shim, Myong Hyon Cho, Michel Kinsy, Tina Wen, G.Edward Suh, and Srinivas Devadas. Static Virtual Channel Allocation in Oblivious Routing. *International Symposium on Networks-on-Chip (NOCS 2009)*, 2009.
- [24] LG Valiant and GJ Brebner. Universal schemes for parallel communication. In *Proceedings of the thirteenth annual ACM symposium on Theory of computing*, pages 263–277. ACM New York, NY, USA, 1981.

- [25] K. Walkowiak. New Algorithms for the Unsplittable Flow Problem. *LECTURE NOTES IN COMPUTER SCIENCE*, 3981:1101, 2006.
- [26] Tina Wen, Myong Hyon Cho, Michel Kinsy, Keun Sup Shim, G.Edward Suh, and Srinivas Devadas. Optimal and Heuristic Bandwidth-Sensitive Oblivious Routing. *The 46th Design Automation Conference (DAC 2009)*, 2009.
- [27] X. Zhong and V.M. Lo. Application-specific deadlock free wormhole routing on multicomputers. In *Proceedings of the 4th International PARLE Conference on Parallel Architectures and Languages Europe*, pages 193–208. Springer-Verlag London, UK, 1992.